






Mean total score (root_cause + local_fix + global_fix, /3)
across 4 synthetic cases — judge=openrouter/openai/gpt-4o




Per-axis means (judge 0/1 per axis, averaged across cases)



Cost per (tier × model) sweep — token-count × OpenRouter pricing
Tier 4 keychain auth = \$0 reported (subscription quota); estimate uses Anthropic published rates



Cost-effectiveness: mean score vs. sweep cost
(4 synthetic cases per (tier × model))



Total tool calls (summed over 4 cases) — agent engagement

